

# CSE 536: Advanced Operating Systems

## Assignment 4: Trap and Emulate Virtualization

**Deadline: 11:59 PM on Dec. 6, 2024**

(100 points)

### 1 Brief Description

In this assignment, you will implement one way to design a virtual machine—*trap and emulate*. The high-level idea of this approach is that the virtual machine (VM) runs as a user-mode process on the host operating system, which also acts like a virtual machine monitor (VMM). The VM is allowed to execute user-level (unprivileged) instructions directly, while all supervisor-level (privileged) instructions *trap* to the VMM. It then *emulates* these trapped privileged instructions and maintains (in-memory) all privileged state of the VM. This state includes all privileged registers (e.g., satp, scause, stvec, etc. in RISC-V).

**GitHub Repo.** Please clone as follows:

```
git clone -b lab4-trap-and-emulate https://github.com/ASTERISC-Teaching/cse536-release.git
```

#### Important note

- **If your assignment does not compile or your upload is corrupted, you will get a zero.** Always double-check your submission.
- **Submit a zip file titled by your ASU username.** For instance, if your username is adil, your file should be adil.zip.

### 2 Description of Tasks

Please follow the steps below to complete this task:

**1: Privileged Virtual Machine State Initialization (10 points).** The VMM must keep in-memory instances of the VM's privileged state. Hence, your first task is to initialize and maintain a data structure that keeps the following privileged registers and VM state:

- Machine trap handling registers
- Machine setup trap registers
- Machine information state registers
- Machine physical memory protection registers
- Supervisor page table register (satp)
- Supervisor trap setup registers
- User trap handling registers
- User trap setup registers
- Current execution mode (e.g., U-mode, S-mode, M-mode).

Please initialize all the registers to 0, except for the `mvendorid` register, which you should initialize to the hexadecimal code of “cse536”. Also, your VM should boot at Machine-mode (M-mode). Define different VM execution modes with different codes decided by you.

### Suggested reading(s)

Please check this website to complete task 1: [RISC-V Privileged Register Mappings](#)

**2. Track Virtual Machine Execution and Redirect Traps (5 points).** All VM processes will start with “vm-”. Hence, your job is to track whether the process is designated as a VM. If it is a VM, the VMM should redirect all user traps that are raised for the execution of privileged instructions to our function `trap_and_emulate()` (defined in `trap-and-emulate.c`). The test cases will be centered around 5 privileged instructions: `csrr`, `csrw`, `sret`, `mret`, and `ecall`. However, note that `ecall` will not be trapped as an error in `usertrap()` function. Make sure you handle the traps for `ecall` correctly too, and redirect its execution.

**3. Decode Trapped Privileged Instruction (20 points).** Once the VM executes a privileged instruction, it will be trapped to the VMM. In this task, your job is to *decode* the instruction so that your VMM can subsequently emulate it. For the decoding process, the VMM must read the instruction and determine (a) the opcode of the instruction (`op`), (b) source registers (`src1` and `src2`), and (c) destination register (`dst`). One of the important parts of this decoding is to map between the codes of registers and the actual source and destination registers. Check the suggested reading #1 for this information.

### Suggested reading(s)

Check this link for the layout of RISC-V instructions: [RISC-V Instruction Layout](#)

**4. Emulate Decoded Instructions (40 points).** Once you have decoded the privileged instruction executed by the VM, you must perform its operation in software. Provided below is a simple (and incomplete) description of what each instruction does:

- `csrr`: Moves the value from a privileged register to an unprivileged register.
- `csrw`: Moves the value from an unprivileged register to a privileged register.
- `sret`: Transfers control from S-mode to U-mode entry point.
- `mret`: Transfers control from M-mode to S-mode entry point.
- `ecall`: Transfers control from U-mode to S-mode or M-mode.  
Also, transfers control from S-mode to M-mode.

Note that all the above mentioned instructions have the same opcode. However, you can differentiate between each instruction by looking at the `src` and `dst` registers. Provided below are some important conditions you should take care of. The list is non-exhaustive; please refer the RISC-V manual.

- An instruction/register should only be accessible in its correct mode inside the VM (e.g., `scause` should not be read in U-mode). If the user process (in the VM) tries to execute a privileged instruction, the VMM should redirect execution to the guest OS’ trap handler.
- Execution of `mret` and `sret` requires that certain registers are set before. For instance, returning from M-mode to S-mode using `mret`, the system writes certain bits of the `mstatus` register, indicating a return to S-mode (check `kernel/start.c`). If the corresponding registers are not correctly set, the execution of these instructions should fail, and the VM should be killed.
- Recall that user and kernel traps are redirected to addresses specified by registers like `stvec`, etc. If you do not remember this, check lectures where we talked about trap handling in `xv6`.
- When the VM writes `0x0` to `mvendorid`, you should shutdown the VM by killing the process. This is of course, not an actual condition, but one we implement to ensure graceful VM shutdown.

**5. Emulating Physical Memory Protection (PMP) (25 points).** There is no paging inside our trap and emulated VMs—all software inside the VM assume that they run directly on physical memory. However, the M-mode software inside the VM can still enable memory protections using PMP. In particular, if the M-mode software writes to the `pmpaddr0` register (e.g., using the instruction `csrw pmpaddr0, s3`) it assumes that part of its physical memory will become inaccessible in U/S-modes (recall assignment #1). Your task is to emulate this for the VM using xv6’s process page tables.

The VMM must perform three sub-tasks to achieve PMP emulation. First, the VMM (xv6) should create a copy of the VM process’ page tables for U/S-mode execution. Second, on the newly-created page tables (which we will call *PMP tables*), the VMM should unmap the regions of memory that are inaccessible (based on the PMP registers). Third, whenever the VM transitions into U/S-mode (e.g., using `sret`), the VMM should switch the VM process’ page tables to PMP tables.

Note that for a created VM, the VMM automatically provides a 4MB memory region from `0x80000000` - `0x80400000`. For the purpose of this assignment, assume that this is the only memory region where PMP will be leveraged for protection. For instance, in a testcase (which you will be provided later), you will be asked to remove permissions for the upper 1MB memory region.

**Testcases.** You are provided a test VM for this assignment (namely `vm-test`). You can change this VM to test different conditions in your code. The output of this VM is provided in `outputs/vm-test`.

## 3 Miscellaneous

### 3.1 Clarifications/Simplifications

Curated below is a list of clarifications:

- In the `mstatus` register, you only need to track the MPP (machine previous privilege) field (ignore all others). Based on this field, your system should return to U/S-mode.
- In the `sstatus` register, you only need to track the SPP (supervisor previous privilege) field (ignore all others). Based on this field, your system should return to U-mode or remain in S-mode.
- You can assume M-mode interrupt/exception delegation registers will always be set to `0xffff` (i.e., all traps are delegated to S-mode).
- You can allow reading of the `mvendorid` register in all modes. Writing to `mvendorid` should only be allowed in M/S-mode.
- For PMP, you will only be asked to remove permissions in a top-of-range (TOR) manner. (Check the example given above and in assignment 1).
- In all testcases, you can assume that all writes to privileged registers will be valid.
- Your system should be able to access the `sie` register, but you can ignore it’s features (i.e., it does nothing).
- If xv6 does not do anything special with a privileged register, our testcases will also not do anything with that register except for read/write from it.

### 3.2 Submitting your Assignment

Please zip the entire provided code directory and submit it to the Canvas under “Assignment 4: CODE”.